

# YASH LALA

[yashlala.com](http://yashlala.com) ◇ [github.com/yashlala](https://github.com/yashlala) ◇ [linkedin.com/in/yashlala](https://www.linkedin.com/in/yashlala)

(510)-400-5572 ◇ [yash.lala@yale.edu](mailto:yash.lala@yale.edu) ◇ Palo Alto, CA

## OVERVIEW

---

I am a PhD student, broadly interested in systems for efficient data-center scale computing. My recent work focuses on far-memory systems.

I have 4+ years of kernel programming expertise, including large-scale projects such as reimplementing the memory-management subsystem in my most recent publication (SOSP '25). I have demonstrated expertise in developing high-throughput, low-latency kernel datapaths — such as a Linux-based system capable of migrating 200 Gib/s of memory at <270 ns per page per core. I'm deeply familiar with CPU microarchitecture and its interactions with the OS, and I excel in principled, data-driven, bottleneck analyses.

In addition to classical OS + networking,

I'm expanding my knowledge to topics other than classical OS+networking. Such as GPU

I am now expanding away from classical OS + networking topics and towards cluster-scale optimizations. I've recently increased scope to cluster-scale optimizations, especially for AI/ML systems. After a summer performance modelling ML network workloads at NVIDIA, I have

## EDUCATION

---

Yale Ph.D. in Computer Science

2023 - Present

UCLA B.S. in Computer Science

GPA: 3.782, 2018 - 2022

## PROFESSIONAL EXPERIENCE

---

### NOVA Lab, Yale CS Department

*PhD Student*

Aug 2023 - Expected 2028

*Advisor: Anurag Khandelwal*

- Broadly researching systems for efficient data-center scale computing. Recent work focuses on far-memory systems.
- Developed techniques for accelerating OS memory migration, aimed towards facilitating high-throughput RDMA-based far memory runtimes. Developed extensive knowledge of the the Linux kernel's memory management, RDMA, page placement, IPI, and swap subsystems. My particular focus was on optimizing high throughput, low-latency runtimes, such as memory migration systems and RDMA SmartNIC network stacks. Publication to appear at **SOSP '25**.
- Developing compact data structures for efficient range queries in high-dimensional spaces. Ongoing project.

### NVIDIA

*Systems Software Intern*

May 2025 - Aug 2025

*Supervisor: Misha Smelyanskiy*

- Developed performance models for Mixture-of-Experts (MoE) communication phases for AI/ML models on NVIDIA's GB200 cluster hardware. Estimated ROI for future communication optimizations, and localized network performance bottlenecks in distributed systems of GPUs.
- Examined MoE token to expert mapping distributions to quantify the workload changes required for optimal hardware utilization. Explored algorithm-network codesign opportunities.
- Profiled existing GPU MoE library implementations (such as NCCL, DeepEP, etc). Optimized GPU MoE libraries for NVIDIA hardware by better utilizing tiered Multi-Node NVLink and Infiniband datacenter networks.
- Gained experience with LLM workload network traffic patterns and GPU performance debugging (particularly in GPU-initiated RDMA network communication).

### SOLAR Lab, UCLA CS Department

*Student Researcher*

Sept 2021 - Sept 2022

*Supervisor: Harry Xu*

- Developed Linux kernel mechanisms for transparent far memory, with the goal of merging my changes upstream. The patchset extends the cpuset controller to allow per-group control of active swap devices. The associated refactoring has positive implications for swap throughput, and makes it easy to manage frontswap-based remote memory systems. Patchset available at [github.com/yashlala/canvas-linux](https://github.com/yashlala/canvas-linux).
- Developed a patchset to improve the Linux kernel's physical page allocation tail latencies by refilling the percpu low-order free page lists asynchronously using RCU.

### CSSI Program, UCLA CS Department

*Tutor Undergrad (TA)*

July 2022

- Taught introductory data science to high school students. Led discussion sections, prepared assignments, graded papers, and advised students.

### Veritas Technologies LLC

*SDE Intern*

June 2021 - Sept 2021

- Worked on NetBackup Flex, a platform for large-scale data consolidation and backup. Implemented automatic node discovery and cluster assimilation over datacenter networks. Primarily worked with Ansible, Docker, and various glue languages.

## PUBLICATIONS

---

### Scalable Far Memory: Balancing Faults and Evictions

August 2023 - August 2025

- Adapted the Linux Kernel to create a far-memory system capable of offloading excess application memory to arbitrary storage backends, for cluster-wide memory utilization improvements. Unlike prior work, this system scales to highly threaded, memory-intensive applications with minimal (near-ideal) application slowdown. Results are applicable to tiered memory systems and all far-memory backends (eg. NVMe, zswap, and RDMA).
- Rearchitected the Linux memory management subsystems to handle high-throughput page faulting and eviction. System achieves >200 Gib/s page migration throughput with only two CPU cores.
- Extensive low-level performance work around the OS-CPU interface, extending well into microarchitectural details. Gained expertise in: x86 IPI overheads and queuing characteristics; IOTLB-related virtualization overheads; TLB and page-walker coherence on x86; RDMA NIC control path throughput bottlenecks due to driver bugs; hypervisor-mediated IPI delivery costs; and lockless programming within the kernel.
- Developed principles for scalable memory migration. Results generalize to a Library-OS based re-implementation.
- Co-first author. Paper to appear at **SOSP 2025**.

### GRU4RecBE: Session Based Recommendations with Features

March 2021 - June 2021

- Developed session-based recommendation system in PyTorch which extends the GRU4REC architecture with rich item features extracted from the pre-trained BERT architecture. Non-attentive model outperforms state-of-the-art session-based models over the MovieLens benchmark datasets. [Paper accepted to AAAI Student track](#).

## PROJECTS

---

### NDN Multicast

Feb 2022 - June 2022

- Extended routing protocols for NDN (Named Data Network) under Lixia Zhang. Extended NLSR (a link-state routing algorithm for NDN) to allow for efficient multicast delivery of NDN Interest packets. Student paper available at [yashlala.com/nlsr-poster.pdf](https://yashlala.com/nlsr-poster.pdf).

## MISCELLANEOUS

---

Enjoys teaching: led student [seminar series](#), volunteered with [ACM TeachLA](#) and as a docent for the [Computer History Museum](#). Certified as a operator for the punched-card based [IBM 1401](#).